

LAN: Learning-based Approximate k -Nearest Neighbor Search in Graph Databases

Yun Peng, Byron Choi, Tsz Nam Chan, Jianliang Xu

Department of Computer Science

Hong Kong Baptist University

{yunpeng, bchoi, edisonchan, xujl}@comp.hkbu.edu.hk

Abstract—The problem of k -nearest neighbor (k -NN) search is fundamental in graph databases, which has numerous real-world applications, such as bioinformatics, computer vision, and software engineering. Graph edit distance (GED) and maximum common subgraph (MCS)-based distance are the most widely used distance measures in k -NN search. However, computing the exact k -NNs of a query graph Q using these measures is prohibitively time-consuming, as a large number of graph distance computations is needed, and computing GED and MCS are both NP-hard. In this paper, we study the *approximate k -nearest neighbor (k -ANN)* search with the aim of trading efficiency with a slight decrease in accuracy. Greedy routing on the proximity graph (PG) index is a state-of-the-art method for k -ANN search. However, such routing algorithms are not designed for graph databases, and simple adoption is inefficient. The core reason is that the exhaustive neighbor exploration at each routing step incurs a large number of distance computations (NDC). In this paper, we propose a learning-based k -ANN search method to reduce NDC. First, we propose to prune unpromising neighbors with distance computations. We use a graph learning model to rank the neighbors at each routing step and explore only the top neighbors. For the accuracy of rank prediction, we propose a neighbor ranking model that works only in the neighborhood of Q . Second, we propose a learning-based method to select the initial node for the routing. The initial node selected has a high probability of being in the neighborhood of Q , such that the neighbor ranking model can be used. Third, we propose a compressed GNN-graph to accelerate the neighbor ranking model and the initial node selection model. We prove that learning efficiency is improved without degrading the accuracy. Our extensive experiments show that our method is about 3.6x to 18.6x faster than the state-of-the-art methods on real-world datasets.

Index Terms—Graph database, Approximate k -NN search, Proximity graph, Learning to route, GNN acceleration

I. INTRODUCTION

The fundamental problem of k -nearest neighbor (k -NN) search in graph databases, which finds the k most similar graphs to a query graph Q , has many applications, such as in cheminformatics [1]–[3], bioinformatics [4], [5], pattern recognition [6]–[8], and software engineering [9]–[11]. For example, in cheminformatics, chemists can use k -NN search to find the molecules with similar structures as the query molecule, as molecules with similar graph structures have similar functions [2]. In software engineering, since the control-flow of a code fragment can be modeled as a graph, software engineers can use k -NN search in a database of control-flow graphs to find code plagiarism issues [10]. Many similarity




Fig. 1: Examples of routing on the proximity graph of a graph database $\mathcal{D} = \{G_0, G_1, \dots, G_{18}\}$ to find the 1-NN G_{17} of Q (The numbers in circles are the distances to Q . The routings are marked by bold arrows. The dashed line marks the ground truth of the neighborhood of Q .)

measures have been proposed to quantify the distance between two graphs, where graph edit distance (GED) and maximum common subgraph (MCS)-based distance are the most widely used measures in k -NN search in graph databases (e.g., [3], [6], [12], [13]). However, computing GED and MCS are both NP-hard [6], [14]. It is impractical to use these measures to find the exact k -NNs. For example, in our preliminary experiments, we use the latest graph similarity search method [15] to find the exact 20-NNs of a randomly selected Q from a dataset AIDS that has 42,687 molecule graphs. This method does not finish after 10 hours. Therefore, this paper considers *k -approximate nearest neighbor (k -ANN)* search in graph databases. To the best of our knowledge, this paper is the first work on k -ANN search in graph databases.

The most state-of-the-art index for k -ANN search to date is the *proximity graph* (PG) [16]–[20]. It supports k -ANN search in a metric space, but it has not previously been studied with graph databases before. The general idea of PG is as follows. Given a database \mathcal{D} , the nodes of PG are the objects in \mathcal{D} and two nodes have an edge if they fulfill a certain proximity property (e.g., navigable small world property [17]). k -ANN search is evaluated through greedy routing on PG. At each routing step, a router computes the distances between the query object Q and all neighbors of the current node of the router and routes to the neighbor that is the closest to Q . The router backtracks if the current node has no neighbor that is closer to Q than itself.

A baseline using PG and existing routing methods for k -ANN search in graph databases is, however, inefficient. The core reason for this inefficiency is that at each routing step,

the router exhaustively computes distances for all neighbors of the current node of the router. Since each node in PG has many neighbors, this kind of exhaustive neighbor exploration incurs a large number of distance computations (NDC), which is incredibly time-consuming, especially considering that computing GED and MCS are both NP-hard. For example, in Fig. 1(a), the routing from G_3 to G_{17} needs to compute distances for all neighbors of G_3 , G_8 , G_{14} , and G_{17} . The NDC is 13, which is 3.25x of the routing length 4. In our preliminary experiments on the AIDS dataset, the NDC can be 20x of the routing length.

In this paper, we propose a learning-based routing method to address k -ANN search in graph databases. Our overall approach is to use graph learning to prune unnecessary distance computations. This approach has three main technical issues, and our solutions can be summarized as follows.

First, we propose routing on PG with neighbor pruning. At each routing step, we prune the neighbors of the current node that are far from the query graph Q from computing the distances. It is motivated by the experimental observation that the routing seldom routes to the neighbors that are far from Q . For example, on the AIDS dataset, at each routing step, we rank the neighbors of the current node by their distances to Q and prune the last 80% of the neighbors; NDC is reduced by 5.7x as a result, while the recall of the search results does not decline. We have proved that *if we have an oracle to rank neighbors by their distances to Q , the routing with neighbor pruning uses a smaller NDC than the baseline, but produces the same search results*. We train a graph learning model M_{rk} to approximate the oracle. At each routing step, suppose G is the current node of the router. We use $M_{rk}(x)$ to predict the top $x\%$ of the neighbors of G . If the prediction of a neighbor G' of G is false, G' is pruned from distance computation. For an efficient training of M_{rk} , we restrict to use M_{rk} only when the router has entered the neighborhood of Q . For the example in Fig. 1(a), the router can use $M_{rk}(30)$ at G_8 to prune G_6 and G_9 from distance computations, as G_6 and G_9 are not among the top 30% of the neighbors of G_8 . Similarly, at G_{14} , G_{13} and G_{15} can be pruned. At G_{17} , G_{18} can be pruned. The NDC is reduced from 13 to 8 as a result. Our experiments show that the efficiency of routing with our learning-based neighbor pruning is ~ 2.9 x higher than that of the baseline on the AIDS dataset.

Second, we propose a learning-based method to select the initial node for the routing. The aim is to ensure that the initial node selected is in the neighborhood \mathcal{N}_Q of Q , s.t. the neighbor pruning model can be used in most routing steps. Specifically, we train a graph learning model M_{nh} to predict if a graph $G \in \mathcal{D}$ is in \mathcal{N}_Q . In the predicted neighborhood $\hat{\mathcal{N}}_Q$, we randomly sample s graphs and compute their distances to Q . The top samples are used as the initial nodes. We have proved that the initial node found has a high probability of being in \mathcal{N}_Q . To reduce the time complexity of initial node selection, we further propose a cluster-based method to prune unpromising graphs in \mathcal{D} . For the example in Fig. 1(a), if G_8 is selected as the initial node instead of G_3 , the NDC is

reduced from 8 to 4. Our experiments show that our initial node selection method can speed up the routing by ~ 2 x when compared with randomly selecting a node as the initial node on the AIDS dataset.

Third, we propose accelerating the graph learning used in the neighbor pruning and initial node selection. Both M_{rk} and M_{nh} take cross-graph learning between a data graph G and a query graph Q as the core module, as it is the latest method for graph distance learning [21], [22]. In the cross-graph learning, a node of G not only aggregates embeddings of its neighbors in G but also pays attention to the embeddings of all nodes in Q . Since many nodes have the same embedding, there are many redundant computations in the cross-graph learning. We propose a compressed GNN-graph (CG) to group together the nodes with the same embedding. We have proved that cross-learning on the CGs of G and Q is more efficient than that on G and Q without degrading the learning accuracy. Our experiments show that the k -ANN search is ~ 1.17 x faster when using our cross-graph learning acceleration on the AIDS dataset. By using these techniques, the 20-ANN queries on the AIDS dataset take ~ 40 seconds on average.

Contributions. The contributions of this paper are as follows.

- We propose a learning-based neighbor pruning method to reduce the number of distance computations in the routing on the PG of a graph database.
- We propose a learning-based initial node selection method. The initial node selected has a high probability of being in the neighborhood of Q .
- We propose a novel compressed GNN-graph (CG). Cross-learning on the CGs of G and Q is equivalent to but faster than that on G and Q .
- Our extensive experiments verify the effectiveness and efficiency of our proposed techniques.

Organizations. The rest of this paper is organized as follows. Sec. II discusses related work. Preliminaries and problem definition are presented in Sec. III. Sec. IV presents the techniques of routing with neighbor pruning. The learning-based initial node selection is presented in Sec. V. Sec. VI presents the graph learning acceleration method. Our experimental evaluation is presented in Sec. VII. Sec. VIII concludes this paper. All proofs are given in Appendix.

II. RELATED WORK

In this section, we summarize existing work that is closely related to this paper.

A. Proximity graphs

Proximity graphs (PGs) are the state-of-the-art indexes for k -ANN search (e.g., [16]–[18], [20], [23]–[29]). Existing works on k -ANN search using PGs can be categorized into two classes. The first class focuses on the proximity properties of PGs. For example, the navigable small world properties are studied in [17] and the relative neighborhood relationships are studied in [18]. The second class focuses on the query evaluations on PGs, which include improving routing efficiency and selecting high-quality initial nodes for routing.

To efficiently route on PGs, Muñoz et al. [30] propose a quadrant-based method to prune unpromising neighbors from distance computations at each routing step. However, this method cannot be used with graph data as graph data have no concept of quadrants. Baranchuk et al. [28] propose a learning-based routing method. However, it is tailor-made for 1-ANN search, and it is not clear how to extend it to efficiently support k -ANN search.

In terms of initial node selection, several works [19], [31] randomly sample a node in a PG as the initial node. DLG [32] and HNSW [17] construct a hierarchy of PGs and use the k -ANN search on the higher level PG to find the initial node for the lower level. However, these methods cannot bound the distance between Q and the initial node selected. Qin et al. [9] propose a learning-based approximate range query method in graph databases. However, if the range threshold is not small, it will take a large NDC, as it can be observed from our experiments.

There are also many indexes for k -ANN search not based on PGs, such as LSH [33] and inverted file index [16]. However, recent studies [20], [34], [35] report that PGs outperform these indexes. Therefore, we omit their details from this subsection.

B. Graph distance learning

Several recent works use graph neural networks (GNNs) to learn the distance between two graphs. A common approach involves first learning the cross-graph embedding of G and Q and then feeding the cross-graph embedding to multilayer perceptrons (MLPs) or convolutional neural networks (CNNs) to predict the distance between G and Q . For example, Li et al. [22] propose a cross-graph attention network GMN. Each node of G aggregates information from both its neighbors in G and all nodes of Q in cross-graph convolution. Bai et al. [36] use the matrix of the inner products of the embeddings of the nodes in G and Q as the cross-graph embedding. Peng et al. [37] propose using a GNN on the association graph of G and Q to learn the cross-graph embedding of the two graphs. A recent work [21] integrates the idea of GMN in the A* search of GED computation. Peng et al. [38] survey the learning-based graph combinatorial optimization methods.

C. Graph neural network acceleration

Many recent works study GNN acceleration. For example, GraphSAGE [39], FastGCN [40], and Adapt [41] sample a subset of neighbors in graph convolution. SGCN [42] removes the non-linear activation functions. Ye and Ji [43] propose a sparse attention technique to accelerate the graph attention network. DegreeQuant [44] and #GNN [45] quantize the node embeddings as integer vectors and binary vectors, respectively. However, these methods do not guarantee that the learning accuracy is preserved after acceleration.

The work that is most closely related to this paper is HAG [46], which accelerates GNN while preserving learning accuracy. HAG aggregates the redundant sum operations in GNN learning. However, HAG cannot reduce the number of matrix multiplications, which is a bottleneck in cross-graph learning.




Fig. 2: An example of two graphs and their GNN-graphs (A and B are node labels)

There are also works that accelerate GNN by taking advantage of the optimized graph primitives of certain platforms (e.g., G^3 [47]) or optimizing GPU kernels (e.g., fuseGNN [48]). These works are orthogonal to this paper, as we focus on reducing redundant computations in GNN learning.

III. PRELIMINARIES AND PROBLEM DEFINITION

This paper studies undirected graphs with node labels. A graph is denoted as $G = (V_G, E_G, \ell_G)$, where V_G and E_G are the node set and the edge set of G , respectively, and $\ell_G(v)$ is the label of node v . The set of neighbors of a node $v \in G$ is denoted by $N_G(v)$. The subscript may be omitted if it is clear from the context.

A. Search semantics

We focus on GED in this paper, as MCS is a special case of GED [49]. A graph G can be transformed into another graph G' through five edit operations: node insertion, edge insertion, node deletion, edge deletion, and node relabeling. The GED of G and G' , denoted by $d(G, G')$, is the smallest number of edit operations that transform G to G' .

Example 1. Fig. 2(a) and (b) show a data graph G and a query graph Q , respectively. $d(G, Q) = 5$.

Approximate k -Nearest Neighbor Search Problem. Given a graph database \mathcal{D} , a query graph Q , and a parameter $k \ll |\mathcal{D}|$, the approximate k -nearest neighbor (k -ANN) search aims to find the k graphs from \mathcal{D} with the smallest distances to Q .

B. Routing on proximity graphs

Given a graph database \mathcal{D} , the nodes of PG are the graphs in \mathcal{D} and two nodes have an edge if they fulfill a certain proximity property (e.g., the navigable small world property [17]). Algorithm 1 presents the logic of a router with which to answer a k -ANN search. Algorithm 1 simply adopts the beam search method of existing PG-based k -ANN search works (e.g., [18]–[20], [27], [31], [34], [35], [50]). Line 1 initializes a priority queue W to store candidates, which is ordered by the distances of the candidates from Q . The binary flag $G.\text{explored}$ marks if the router has computed distances for the neighbors of G . Line 3 selects the initial node for the routing. Lines 5–10 perform the routing steps. At each routing step, the router picks the unexplored node in W with the smallest distance to Q as the current node. When there is a tie, the node ID is

Algorithm 1 Greedy routing on proximity graph (baseline)

Input: PG \mathcal{G} , query Q , beam size b , parameter k
Output: k -ANNs of Q

- 1: initialize a priority queue $W = \emptyset$ as the pool to store candidates
- 2: G_{explored} is **false** by default for each node G in \mathcal{G}
- 3: $G_{\text{init}} = \text{select an initial node in } \mathcal{G}$
- 4: $W.\text{add}((d(G_{\text{init}}, Q), G_{\text{init}}))$ $\triangleright W$ is in ascending order of d
- 5: **while** W has unexplored nodes **do**
- 6: $G = \text{the unexplored node in } W \text{ with the smallest GED to } Q$
- 7: $\text{neigh_explore}(\mathcal{G}, G, Q)$
- 8: $G.\text{explored} = \text{true}$
- 9: resize W to size b
- 10: **end while**
- 11: $\mathcal{R}_{bs} = \text{top-}k \text{ in } W$
- 12: **return** \mathcal{R}_{bs}
- 13: **function** $\text{neigh_explore}(\mathcal{G}, G, Q)$
- 14: **for** each neighbor G' of G in \mathcal{G} **do**
- 15: add $(d(G', Q), G')$ into W
- 16: **end for**

used to break the tie. Then, the router computes distances for all neighbors of the current node and adds all of them to W (Line 7). Then, W is resized to b (Line 9). During the resizing of W , for two nodes in W that have the same distance to Q , the unexplored one has a higher priority; if both are explored, the most recently explored one has a higher priority; and if both are unexplored, the one with a smaller node ID has a higher priority. The routing stops if all nodes in W have been explored and the top- k graphs in W are returned.

C. Graph neural network

The main idea of Graph neural networks (GNNs) is to use graph convolutions to learn the embeddings of graphs. In this paper, we adopt the well-known GIN model [51]. Given a graph G , the graph convolution of GIN is as follows.¹ For a node u of G ,

$$\mathbf{h}_u^l = \text{ReLU}(\mathbf{W}^l(\mathbf{h}_u^{l-1} + \sum_{v \in N(u)} \mathbf{h}_v^{l-1})) \quad (1)$$

where l denotes the layer ID, \mathbf{h}_u^l is the embedding of u at the l -th layer, v is a neighbor of u , and \mathbf{W}^l is the trainable parameter matrix, respectively. \mathbf{h}_u^0 is the input feature of u (e.g., the one-hot encoding of $\ell(u)$). If there are L layers, the embedding of G is $\mathbf{h}_G = \text{mean}_{u \in G} \mathbf{h}_u^L$.

The work [51] has proved that GIN is equivalent to the Weisfeiler-Lehman (WL) labeling, which is a well-known technique used to test graph isomorphism. Given a graph $G = (V, E, \ell)$, WL labeling iteratively computes the WL labels of the nodes in G . The WL label of a node $u \in G$ at the l -th iteration, denoted by $wl^l(u)$, is as below.

$$wl^l(u) = wl^{l-1}(u), \{\{wl^{l-1}(v) | v \in N(u)\}\}, \text{ and} \quad (2)$$

$$wl^0(u) = \ell(u), \quad (3)$$

where $\{\{\}\}$ denotes a multi-set.

¹The coefficient $1 - \epsilon$ of \mathbf{h}_u^{l-1} is omitted for the sake of simplicity, which does not affect the results of this paper.

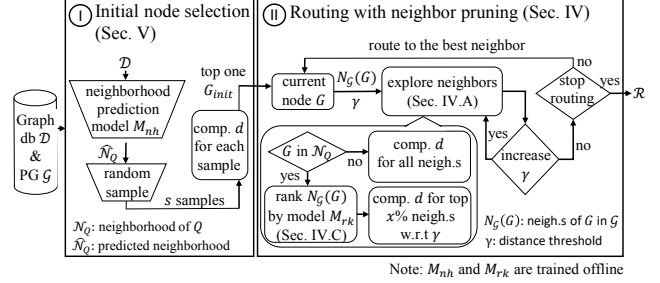


Fig. 3: Overview of our learning-based k -ANN search

The nodes in G with the same WL label at the l -th iteration of WL labeling must have the same embedding at the l -th layer of GIN.

D. GNN-graph

The work [46] represents the computation of graph convolutions of a GNN using a GNN-graph. Given a graph G and a GNN with L graph convolution layers, the GNN-graph of G is an $L + 1$ -level directed acyclic graph $H_{G,L} = (V_0, V_1, \dots, V_L, E)$, where $V_l = \{\mathbf{h}_u^l | u \in G\}$.² For two nodes \mathbf{h}_u^l and \mathbf{h}_v^{l-1} , $(\mathbf{h}_v^{l-1}, \mathbf{h}_u^l) \in H_{G,L}$ if $(u, v) \in G$. For each $u \in G$, $(\mathbf{h}_u^{l-1}, \mathbf{h}_u^l)$ is in $H_{G,L}$. Since many nodes of G have the same label, many nodes in $H_{G,L}$ have identical embeddings.

Example 2. For the graphs G and Q in Fig. 2(a) and (b), Fig. 2(c) and (d) shows the GNN-graphs $H_{G,2}$ and $H_{Q,2}$, respectively. Since $\ell(v_1) = \ell(v_2) = \ell(v_3)$ and $\ell(u_0) = \ell(u_2)$, $\mathbf{h}_{v_1}^l = \mathbf{h}_{v_2}^l = \mathbf{h}_{v_3}^l$ and $\mathbf{h}_{u_0}^l = \mathbf{h}_{u_2}^l$, for $l = 0, 1, 2$.

E. Cross-graph learning

Cross-graph learning is the state-of-the-art technique for graph distance learning [21], [22]. Given two graphs G and Q , each node u of G not only aggregates embeddings from its neighbors in G but also pays attention to the embeddings of all nodes in Q . The cross-graph learning is defined as below.

Definition 1. Given two graphs G and Q , for each node $u \in G$, the embedding of u at the l -th layer, $l \geq 1$, is computed as follows.

$$\mathbf{h}_u^l = \text{ReLU}(\mathbf{W}^l(\mathbf{h}_u^{l-1} + (\sum_{u' \in N(u)} \mathbf{h}_{u'}^{l-1}) + \boldsymbol{\mu}_u^l)), \quad (4)$$

$$\boldsymbol{\mu}_u^l = \sum_{v \in Q} \alpha_{u,v} \mathbf{h}_v^{l-1}, \text{ and} \quad (5)$$

$$\alpha_{u,v} = \frac{\exp(\mathbf{a} \cdot (\mathbf{h}_u^{l-1} \parallel \mathbf{h}_v^{l-1}))}{\sum_{v' \in Q} \exp(\mathbf{a} \cdot (\mathbf{h}_u^{l-1} \parallel \mathbf{h}_{v'}^{l-1}))}, \quad (6)$$

where \parallel denotes concatenation, \mathbf{a} is a trainable parameter vector, \cdot denotes the inner product, and α is attention weight. \mathbf{h}_u^0 is the input feature of u (e.g., the one-hot encoding of $\ell(u)$). If there are L layers, $\mathbf{h}_G = \text{mean}_{u \in G} \mathbf{h}_u^L$. The cross-graph embedding of G and Q is $\mathbf{h}_{G,Q} = \mathbf{h}_G \parallel \mathbf{h}_Q$.

²We slightly abuse the symbol \mathbf{h}_u^l to denote an embedding vector and a node in V_l of H .

Algorithm 2 Routing with neighbor pruning (np_route)

Input: PG \mathcal{G} , query Q , beam size b , answer count k , oracle \mathcal{O} , step size d_s
Output: k -ANNs of Q

```

1: initialize a priority queue  $W = \emptyset$  as the pool to store candidates
2:  $G_{\text{explored}}$  is false by default for each node  $G$  in  $\mathcal{G}$ 
3:  $G_{\text{init}} = \text{select a node in } \mathcal{G} \text{ as the initial node}$ 
4:  $W.\text{add}((d(G_{\text{init}}, Q), G_{\text{init}}))$   $\triangleright W$  is in ascending order of  $d$ 
5:  $G = \text{the node in } W \text{ with the smallest GED to } Q$ 
6: while  $G_{\text{explored}} = \text{false}$  do
7:    $\text{rank\_expl}(\mathcal{G}, G, Q, W, d(G, Q), \mathcal{O})$ 
8:    $G_{\text{explored}} = \text{true}$ 
9:    $\text{resize } W \text{ to } b$ 
10:   $G = \text{the node in } W \text{ with the smallest } d$ 
11: end while
12:  $G_{\text{flo}} = \text{the node in } W \text{ with the smallest } d$   $\triangleright$  first local optimal
13:  $\gamma = d(G_{\text{flo}}, Q) + d_s$ 
14: while true do
15:   for each  $G$  that is explored do
16:      $\text{all\_quali\_neigh}(\mathcal{G}, G, Q, W, \gamma, \mathcal{O})$ 
17:   end for
18:    $\text{resize } W \text{ to } b$ 
19:   if all nodes in  $W$  have been explored then
20:     break  $\triangleright$  stop condition of routing
21:   end if
22:   while  $W$  has unexplored nodes with  $d \leq \gamma$  do
23:      $G = \text{the unexplored node in } W \text{ with the smallest } d$ 
24:      $\text{rank\_expl}(\mathcal{G}, G, Q, W, \gamma, \mathcal{O})$ 
25:      $G_{\text{explored}} = \text{true}$ 
26:      $\text{resize } W \text{ to } b$ 
27:   end while
28:    $\gamma = \gamma + d_s$ 
29: end while
30:  $\mathcal{R}_{np} = \text{top-}k \text{ in } W$ 
31: return  $\mathcal{R}_{np}$ 

```

F. Solution overview

Fig. 3 shows the overview of our learning-based k -ANN search. Given a graph database \mathcal{D} and a PG \mathcal{G} of \mathcal{D} , we propose a learning-based routing on \mathcal{G} to efficiently find the k -ANNs of Q . The key is using graph learning models to prune unnecessary distance computations. Given a query graph Q , we maintain a priority queue W to store the candidates. ① We use a model M_{nh} to predict the neighborhood \mathcal{N}_Q of Q . From the predicted neighborhood $\hat{\mathcal{N}}_Q$, we randomly sample s graphs and compute their distances from Q . The top sample is used as the initial node G_{init} for the routing on \mathcal{G} and G_{init} is added to W . ② At each routing step, let G be the current node of the router. If G is not in \mathcal{N}_Q , we simply compute distances for all neighbors of G . If G is in \mathcal{N}_Q , we use a model M_{rk} to rank the neighbors of G and only compute distances for the top $x\%$ of neighbors. The $x\%$ is controlled by a distance threshold γ , which can be updated as the routing progresses, if needed. The neighbors whose distances from Q are computed are added to W . If the stopping condition is met, the routing stops and the top- k of W is returned as the search results \mathcal{R} . The construction of \mathcal{G} and training of the models M_{nh} and M_{rk} take place offline.

IV. ROUTING WITH NEIGHBOR PRUNING

The main idea of routing with neighbor pruning is that when the router reaches a node G in the PG, we do not compute distances for all neighbors of G at once. We rank the neighbors and only compute distances for the top $x\%$ neighbors of G .

Algorithm 3 Get all qualified neighbors (all_quali_neigh)

Input: PG \mathcal{G} , node G of \mathcal{G} , query Q , pool W , threshold γ , oracle \mathcal{O}

```

1: suppose  $G$ 's neighbors are ranked by  $\mathcal{O}$  and partitioned to  $B_0, B_1, \dots, B_n$ 
2: suppose  $B_0, B_1, \dots, B_i$  have been opened
3: for  $j = 0$  to  $i$  do
4:   for each  $G'$  in  $B_j$  that is unexplored do
5:      $\text{add } (d(G', Q), G') \text{ into } W$ 
6:   end for
7:   if Line 5 adds a neighbor  $G'$  in  $B_j$  to  $W$  and  $d(Q, G') \geq \gamma$  then
8:     return
9:   end if
10: end for
11: for  $j = i + 1$  to  $n$  do
12:   for each neighbor  $G'$  in  $B_j$  do  $\triangleright$  open the batch  $B_j$ 
13:      $\text{compute } d(G', Q) \text{ and add } (d(G', Q), G') \text{ into } W$ 
14:   end for
15:   if Line 13 adds a neighbor  $G'$  in  $B_j$  to  $W$  and  $d(Q, G') \geq \gamma$  then
16:     return
17:   end if
18: end for
19: return

```

Algorithm 4 Neighbor ranking and exploration (rank_expl)

Input: PG \mathcal{G} , node G of \mathcal{G} , query Q , pool W , threshold γ , oracle \mathcal{O}

```

1: suppose  $G$ 's neighbors are ranked by  $\mathcal{O}$  and partitioned to  $B_0, B_1, \dots, B_n$ 
2: suppose  $B_0, B_1, \dots, B_i$  have been opened
3:  $G'$  is the neighbor farthest from  $Q$  in the opened batches of  $G$ 's neighbors
4: if  $d(Q, G') \geq \gamma$  then
5:   return
6: end if
7: for  $j = i + 1$  to  $n$  do
8:   for each neighbor  $G'$  in  $B_j$  do  $\triangleright$  open the batch  $B_j$ 
9:      $\text{compute } d(G', Q) \text{ and add } (d(G', Q), G') \text{ into } W$ 
10:   end for
11:   if  $B_j$  contains a neighbor  $G'$  satisfying  $d(Q, G') \geq \gamma$  then
12:     return
13:   end if
14: end for

```

The $x\%$ is controlled by a distance threshold, which can be updated as the routing progresses.

In the following, we first present the algorithm of routing with neighbor pruning under the assumption of an oracle for neighbor ranking. Then, we remove the assumption by proposing a graph learning model to approximate the oracle.

A. Algorithm of routing with neighbor pruning

Assume we have an oracle that can rank the neighbors of G in a negligible amount of time. Then, we can use the oracle to partition the neighbors of G into batches B_0, B_1, \dots, B_n of equal size in order to control the percentage of neighbors in distance computation. Each batch has $y\%$ neighbors of G , where y is a tunable parameter, and all neighbors in B_i are farther from Q than all neighbors in B_{i-1} , for $i = 1, \dots, n$. Given a distance threshold γ , we sequentially open the batches B_0, B_1, \dots, B_i and compute distances for the neighbors in the opened batches. If B_i has a neighbor G' satisfying $d(G', Q) > \gamma$, we do not open the batches after B_i . We call G explored (*i.e.*, $G_{\text{explored}} = \text{true}$) if at least one batch of the neighbors of G is opened.

np_route (Algorithm 2) presents the algorithm of routing with neighbor pruning using the oracle. The routing has two stages. The first stage is routing without backtracking until the

first local optimal G_{flo} is reached (Lines 1-12). The second stage is routing with backtracking (Lines 13-29).

In the first stage, at each routing step, let G be the current node of the router, $d(Q, G)$ is used as the GED threshold to control the percentage of neighbor exploration at G (Line 7).

The second stage has several while-loops (Lines 14-29). In each while-loop (Lines 14-29), there is a GED threshold γ , and the routing is restricted in the nodes with GEDs to Q that do not exceed γ (Lines 22-27). To avoid missing a qualified node, for each explored node G , Lines 15-17 add all the unexplored neighbors of G with distances to Q that are no more than γ into W (note that G may not be in W). If W has no unexplored node within distance γ from Q , γ is increased by d_s (Line 28). The for-loop (Lines 15-17) adds all qualified neighbors to W with regard to the new γ . The routing stops if all nodes in W are explored (Line 19), and the top- k in W is returned.

Example 3. Consider the example shown in Fig. 1(b). Suppose $y = 30$, $b = 2$, $k = 1$, and $d_s = 1$. The first stage involves routing from the initial node G_2 to G_7 and the second stage involves routing from G_7 to G_{17} .

The first stage: At G_2 , $d(Q, G_2) = 7$ is used as the distance threshold for neighbor exploration. Since $y = 30$, the batches $B_{0,G_2} = \{G_6\}$ and $B_{1,G_2} = \{G_3\}$ of G_2 's neighbors are opened. $W = [G_6, G_3]$. Since G_6 is the best in W , the router explores G_6 . At G_6 , $d(Q, G_6) = 4$ is used as the distance threshold. The batches $B_{0,G_6} = \{G_7\}$ and $B_{1,G_6} = \{G_8\}$ of G_6 's neighbors are opened. $W = [G_7, G_8]$. Since G_7 is the best in W , the router explores G_7 . At G_7 , $d(Q, G_7) = 2$ is used as the distance threshold. The batch $B_{0,G_7} = \{G_{13}\}$ of G_7 's neighbors is opened. $W = [G_7, G_{13}]$. Since $d(Q, G_{13}) > d(Q, G_7)$, G_7 is the first local optimal.

The second stage: Since $d_s = 1$, $\gamma = d(Q, G_7) + d_s = 3$, which is used as the distance threshold (Line 13). The for-loop (Lines 15-17) makes sure that no qualified neighbor of the explored nodes is missed. Line 16 for G_2 adds G_3 into W and Line 16 for G_6 adds G_8 into W . At this point, $W = [G_7, G_{13}, G_8, G_3]$. Line 18 updates W as $[G_7, G_{13}]$. Since G_{13} is unexplored and $d(Q, G_{13}) \leq \gamma = 3$, Line 24 explores the neighbors of G_{13} . The batch $B_{0,G_{13}} = \{G_7, G_{14}\}$ of G_{13} 's neighbors is opened. $W = [G_{14}, G_7, G_{13}]$. Line 26 updates W as $[G_{14}, G_7]$. Then, Line 24 explores G_{14} . Since $\gamma = 3$, the batches $B_{0,G_{14}} = \{G_{17}\}$ and $B_{1,G_{14}} = \{G_{16}, G_{13}\}$ of G_{14} 's neighbors are opened. At this point $W = [G_{17}, G_{16}, G_{14}, G_7, G_{13}]$. Line 26 updates W as $[G_{17}, G_{16}]$. Then, Line 24 explores G_{17} . At G_{17} , since $\gamma = 3$, all batches of G_{17} 's neighbors are opened and $W = [G_{17}, G_{16}, G_{14}, G_{18}]$. Line 26 updates W as $[G_{17}, G_{16}]$. Since $d(Q, G_{16}) < \gamma = 3$, Line 24 explores G_{16} . Since $\gamma = 3$, all batches of G_{16} 's neighbors are opened and $W = [G_{17}, G_{16}, G_{14}, G_{15}]$. Line 26 updates W as $[G_{17}, G_{16}]$. Since W has no unexplored node, Line 28 increases $\gamma = 3 + d_s = 4$. The for-loop (Lines 15-17) adds G_3 , G_8 , G_{12} , G_{15} , and G_{18} into W . At this point $W = [G_{17}, G_{16}, G_{18}, G_8, G_{15}, G_{12}, G_3]$. Line 18 updates W as $[G_{17}, G_{16}]$. Since all nodes in W have been explored,

Line 19 stops the routing and G_{17} is returned.

B. Analysis of np_route with an oracle neighbor ranker

We analyze the performance of np_route using the oracle neighbor ranker (Algorithm 2) by comparing the search results and the NDC with baseline (Algorithm 1).

Lemma 1. Given the same initial node G_{init} and beam size b , the sequence of nodes explored by np_route is the same as that of baseline, where for a node G , the time when G is explored is the time when $G.explored$ is set to true.

Theorem 1. Let \mathcal{R}_{np} and \mathcal{R}_{bs} denote the sets of graphs returned by np_route and baseline, respectively. Given the same initial node G_{init} and beam size b , $\mathcal{R}_{np} = \mathcal{R}_{bs}$. The NDC used by np_route is no more than that of baseline. The space cost of np_route is at most two times of that of baseline.

As presented in a recent survey [50] there are some variants of baseline. Their space complexities are the same as that of np_route in the worst case.

C. Learning-based neighbor ranking

One may attempt to directly train a model to rank the neighbors of G . However, it is technically challenging to obtain such a model to fully rank the neighbors of G . Since each batch has $y\%$ of G 's neighbors, we propose training $100/y$ binary rankers and the i -th ranker M_{rk}^i partially ranks the neighbors of G . Specifically, M_{rk}^i classifies G 's neighbors into two classes: the positive class consisting of the top $iy\%$ neighbors and the negative class consisting of the remaining neighbors.

1) *Model design:* For a node G and a query Q , for each neighbor G' of G , M_{rk}^i first learns the cross-graph embedding $\mathbf{h}_{G',Q}$ of G' and Q (Sec. III-E) and then feeds the concatenation of $\mathbf{h}_{G',Q}$ and \mathbf{h}_G into a multilayer perceptron (MLP) to make a binary classification. The binary cross-entropy with the regularizer in [37] is used as the loss function.

2) *Model training:* Assume we have a query workload \mathcal{Q} , which can be historical queries or be sampled from the database \mathcal{D} [9]. A basic method to train M_{rk}^i is that for each $Q \in \mathcal{Q}$, we use each node G in \mathcal{G} and G 's neighbors as the training data. Formally, the training data are $\{(Q, G', G) | Q \in \mathcal{Q}, G' \in N_G(G), G \in \mathcal{G}\}$. The class label of (Q, G', G) is positive if G' is among the top $iy\%$ of G 's neighbors; otherwise, it is negative. However, the training data can be huge. It may take a long time to train M_{rk}^i .

Recent studies [26], [35] find that most routing steps are in the neighborhood \mathcal{N}_Q of Q . \mathcal{N}_Q can be defined as $\{G | d(Q, G) < \gamma^*, G \in \mathcal{G}\}$, where γ^* is a tunable parameter. Motivated by this observation, we propose using M_{rk}^i only after the router enters the neighborhood of Q . This means that we only need to use $\{(Q, G', G) | Q \in \mathcal{Q}, G' \in N_G(G), G \in \mathcal{N}_Q\}$ as the training data to train M_{rk}^i . Since $|\mathcal{N}_Q| \ll |\mathcal{G}|$, the size of training data is significantly reduced, and hence M_{rk}^i can be efficiently trained.

3) *Routing with learning-based neighbor ranking*: A flow chart of the routing with learning-based neighbor ranking is shown in Fig. 3. Specifically, in Algorithms 3 and 4, suppose we need to open the batch B_j of the neighbors of G . If $G \notin \mathcal{N}_Q$, we directly compute $d(G', Q)$ for each neighbor G' of G ; otherwise, we use the j -th model M_{rk}^j to make a prediction for each neighbor G' of G , and compute $d(G', Q)$ only for the neighbors with positive predictions. For each prediction, if we use the existing cross-graph learning method (Def. 1), the time complexity is $O(|V_Q||V_{G'}|)$. In Sec. VI, we will propose a cross-graph learning acceleration technique.

V. LEARNING-BASED INITIAL NODE SELECTION

Since the neighbor ranking models work only in the neighborhood \mathcal{N}_Q of Q , the initial node of the routing should be in \mathcal{N}_Q in order to use the neighbor pruning technique in most routing steps. In the following, we first present the initial node selection method given a neighborhood prediction model M_{nh} . Then, we present the design of M_{nh} .

A. Initial node selection method

Suppose we have trained a model M_{nh} using the query workload \mathcal{Q} , such that for each graph G in \mathcal{D} , M_{nh} can predict if G is in the neighborhood of Q . Then, we randomly sample s graphs in the predicted neighborhood $\hat{\mathcal{N}}_Q$ and compute the GEDs to Q for the samples. The samples with the smallest GEDs to Q are used as the initial nodes. This method has a high probability of finding a graph in \mathcal{N}_Q .

Lemma 2. Suppose the precision of the prediction of M_{nh} is p (i.e., $p = |\hat{\mathcal{N}}_Q \cap \mathcal{N}_Q|/|\hat{\mathcal{N}}_Q|$). The probability that the s samples have at least a graph in \mathcal{N}_Q is $1 - (1 - p)^s$.

B. Neighborhood prediction model

We first present a basic design of the neighborhood prediction model. Then, we propose an optimized design.

1) *Basic design of M_{nh}* : Given a data graph G and a query Q , M_{nh} first learns the cross-graph embedding $\mathbf{h}_{G,Q}$ (Sec. III-E) and then feeds $\mathbf{h}_{G,Q}$ to an MLP to make a binary prediction. If G is in \mathcal{N}_Q , it is positive; otherwise, it is negative. The binary cross entropy with the regularizer [37] is used as the loss function. Since the number of negative samples is much larger than that of positive samples, the negative class downsampling technique [52] is used to train M_{nh} .

This design of M_{nh} has a shortcoming. To predict \mathcal{N}_Q , we need to make a prediction for each graph in \mathcal{D} . The number of predictions is $O(|\mathcal{D}|)$, which is inefficient when \mathcal{D} is large.

2) *Optimized design of M_{nh}* : We adopt the cluster-based learning framework [53] to reduce the number of predictions. The main idea is to cluster the graphs in \mathcal{D} into a set of clusters \mathcal{C} . We train a model M_c to predict the size of the intersection between each cluster C in \mathcal{C} and \mathcal{N}_Q for each query Q in \mathcal{Q} . The distribution of the intersection size between \mathcal{N}_Q and the clusters is skewed. Hence, we use a neural network to learn the distribution.

The number of predictions is reduced by using M_c . Specifically, for a query Q , we use M_c to make a prediction for each cluster. The top clusters are selected. For each selected cluster C , we use M_{nh} to make a prediction for each graph in C . The total number of predictions is reduced to $|\mathcal{C}| + \sum_{C \in \mathcal{C}'} |C|$, where \mathcal{C}' denotes the selected clusters.

To cluster the graphs in \mathcal{D} , we can use existing graph embedding techniques (e.g., [54]) to compute the embeddings of the graphs in \mathcal{D} and then use KMeans for clustering.

VI. CROSS-GRAPH LEARNING ACCELERATION BASED ON GNN-GRAPH COMPRESSION

The cross-graph learning in Definition 1 has many redundant computations because many nodes in a graph have identical embedding (see Example 2). To eliminate the redundancy, in this section, we first propose a compressed GNN-graph (CG), where the nodes with the same embedding are grouped together. Then, we present the cross-graph learning method using two CGs. After that, we present the optimum CG construction method.

A. Compressed GNN-graph

Definition 2. Given an $L + 1$ -level GNN-graph $H_{G,L} = (V_0, V_1, \dots, V_L, E)$ of a graph G , the compressed GNN-graph (CG) of G is an edge weighted $L + 1$ -level directed acyclic graph $H_{G,L}^* = (V_0, V_1, \dots, V_L, E, w)$.

- For the l -th level, $l = 0, \dots, L$, suppose the nodes in $V_l(H_{G,L})$ can be grouped into a set of groups, such that for any two nodes \mathbf{h}_u^l and \mathbf{h}_v^l in a group, the embedding vectors \mathbf{h}_u^l and \mathbf{h}_v^l are equal. $V_l(H_{G,L}^*)$ has a node for each group. $|g|$ denotes the size of a group g .³
- For two nodes $g_{l-1,i} \in V_{l-1}(H_{G,L}^*)$ and $g_{l,j} \in V_l(H_{G,L}^*)$, $(g_{l-1,i}, g_{l,j}) \in E(H_{G,L}^*)$ if $H_{G,L}$ has edges crossing $g_{l-1,i}$ and $g_{l,j}$.
- Let \mathbf{h}_u^l be any node in the group $g_{l,j}$. The weight w of $(g_{l-1,i}, g_{l,j})$ is the number of incoming neighbors of \mathbf{h}_u^l in the group $g_{l-1,i}$ in $H_{G,L}$.

For presentation simplicity, in the following, we use g and q to denote the nodes in the CGs of G and Q , respectively.

Example 4. Fig. 4(a) presents the CG $H_{G,2}^*$ of the GNN-graph $H_{G,2}$ in Fig. 2(c). Since the embedding vectors $\mathbf{h}_{v_1}^0 = \mathbf{h}_{v_2}^0 = \mathbf{h}_{v_3}^0$, the nodes of $V_0(H_{G,2})$ are grouped into two groups $g_{0,0} = \{\mathbf{h}_{v_1}^0\}$ and $g_{0,1} = \{\mathbf{h}_{v_2}^0, \mathbf{h}_{v_3}^0\}$. Hence, $V_0(H_{G,2}^*)$ has two nodes $g_{0,0}$ and $g_{0,1}$. Similarly, $V_1(H_{G,2}^*)$ has two nodes $g_{1,0}$ and $g_{1,1}$, and $V_2(H_{G,2}^*)$ has two nodes $g_{2,0}$ and $g_{2,1}$. The weighted edges are added based on Definition 2. For instance, $\mathbf{h}_{v_1}^1$ is in the group $g_{1,0}$. $\mathbf{h}_{v_0}^1$ has one and three incoming neighbors in the groups $g_{0,0}$ and $g_{0,1}$ in $H_{G,2}$, respectively. Hence, $w(g_{0,0}, g_{1,0}) = 1$ and $w(g_{0,1}, g_{1,0}) = 3$. For the GNN-graph $H_{Q,2}$ in Fig. 2(d), the CG $H_{Q,2}^*$ is shown in Fig. 4(b).

³We slightly abuse the symbol g to denote a node in $H_{G,L}^*$ and a group of nodes in $H_{G,L}$.




Fig. 4: CGs of G and Q in Fig. 2 (edge weight = 1 is omitted for clarity)

\mathbf{h}^0 affects the compression. The more nodes with the same \mathbf{h}^0 , the better. In real datasets, many nodes have the same node label. For instance, on the AIDS dataset with 42,687 graphs, a graph has 25 nodes but only has 4 distinct node labels on average. If the one-hot encoding of the node label is used to compute \mathbf{h}^0 , which is common in practice, many nodes have the same \mathbf{h}^0 .

B. Cross-graph learning on CGs

We define the cross-graph learning on the CGs of G and Q as follows.

Definition 3. Given the CGs $H_{G,L}^*$ and $H_{Q,L}^*$ of G and Q , respectively, cross-graph learning computes the embedding $\mathbf{h}_{g_{l,i}}$ of each node $g_{l,i}$ in the l -th level of $H_{G,L}^*$, $l = 1, \dots, L$, as follows:

$$\mathbf{h}_{g_{l,i}} = \text{ReLU}(\mathbf{W}^l(\mathbf{t}_{g_{l,i}} + \boldsymbol{\mu}_{g_{l,i}})), \quad (7)$$

$$\mathbf{t}_{g_{l,i}} = \sum_{g \in N_{H_{G,L}^*}^{in}(g_{l,i})} w(g, g_{l,i}) \mathbf{h}_g, \quad (8)$$

$$\boldsymbol{\mu}_{g_{l,i}} = \sum_{q \in V_l(H_{Q,L}^*)} |q| \alpha_{g_{l,i},q} \mathbf{h}_q, \text{ and} \quad (9)$$

$$\alpha_{g_{l,i},q} = \frac{\exp(\mathbf{a} \cdot (\mathbf{t}_{g_{l,i}} \parallel \mathbf{h}_q))}{\sum_{q' \in V_l(H_{Q,L}^*)} |q'| \exp(\mathbf{a} \cdot (\mathbf{t}_{g_{l,i}} \parallel \mathbf{h}_{q'}))}, \quad (10)$$

where N^{in} denotes the incoming neighbors. For each node $g_{0,i}$ in the 0-th level of $H_{G,L}^*$, $\mathbf{h}_{g_{0,i}} = \mathbf{h}_u^0$ for any \mathbf{h}_u^0 in the group $g_{0,i}$. The embedding of $H_{G,L}^*$ is the weighted average of the embeddings of the nodes in the L -th level of $H_{G,L}^*$, i.e., $\mathbf{h}_{H_{G,L}^*} = (\sum_{g \in V_L(H_{G,L}^*)} |g| \mathbf{h}_g) / (\sum_{g \in V_L(H_{G,L}^*)} |g|)$. The cross-graph embedding of G and Q is $\mathbf{h}_{H_{G,L}^*} \parallel \mathbf{h}_{H_{Q,L}^*}$.

Example 5. Continue with Example 4, $\mathbf{t}_{g_{1,0}} = \mathbf{h}_{g_{0,0}} + 3 \times \mathbf{h}_{g_{0,1}}$. $\boldsymbol{\mu}_{g_{1,0}} = 2 \times \alpha_{g_{1,0},q_{1,0}} \mathbf{h}_{q_{1,0}} + \alpha_{g_{1,0},q_{1,1}} \mathbf{h}_{q_{1,1}}$. $\mathbf{h}_{g_{1,0}} = \text{ReLU}(\mathbf{W}^1(\mathbf{t}_{g_{1,0}} + \boldsymbol{\mu}_{g_{1,0}}))$. $\mathbf{h}_{H_{G,2}^*} = (\mathbf{h}_{g_{2,0}} + 3 \times \mathbf{h}_{g_{2,1}}) / 4$. Similarly, $\mathbf{h}_{H_{Q,2}^*} = (2 \times \mathbf{h}_{q_{2,0}} + \mathbf{h}_{q_{2,1}}) / 3$. $\mathbf{h}_{H_{G,2}^*} \parallel \mathbf{h}_{H_{Q,2}^*}$ is the cross-graph embedding.

Theorem 2. Given two graphs G and Q , $\mathbf{h}_{H_{G,L}^*} \parallel \mathbf{h}_{H_{Q,L}^*}$ computed by the cross-graph learning on the CGs $H_{G,L}^*$ and $H_{Q,L}^*$ (Definition 3) equals to $\mathbf{h}_G \parallel \mathbf{h}_Q$ computed by the cross-graph learning on G and Q (Definition 1).

Algorithm 5 CG construction

Input: graph G and layer number L
Output: the CG $H_{G,L}^*$ of G
1: initialize $H_{G,L}^* = (V_0, V_1, \dots, V_L, E, w)$, $V_l = \emptyset$, $E = \emptyset$
2: perform L iterations of WL labeling
3: **for** each $l = 0$ to L **do**
4: group the nodes of G by w^l and V_l has a node for each group
5: **end for**
6: **for** each node $g_{l,j} \in V_l$ and each node $g_{l-1,i} \in V_{l-1}$ **do**
7: let u be a node of G in the group $g_{l,j}$
8: the weight of the edge $w(g_{l-1,i}, g_{l,j}) = |N_G(u) \cap g_{l-1,i}|$
9: $w(g_{l-1,i}, g_{l,j}) += 1$ if u is also in the group $g_{l-1,i}$
10: **end for**
11: **return** $H_{G,L}^*$

Dataset	#graphs	avg $ V $	avg $ E $	#nlabel
AIDS	42,687	25.6	27.5	51
LINUX	47,239	35.5	37.7	36
PUBCHEM	22,794	48.2	50.8	10
SYN	1,000,000	10.1	15.9	5

TABLE I: Statistics of datasets

Theorem 3. The time complexity of cross-graph learning on two CGs $H_{G,L}^*$ and $H_{Q,L}^*$ is $O((|V(H_{G,L}^*)| + |V(H_{Q,L}^*)|) + (|E(H_{G,L}^*)| + |E(H_{Q,L}^*)|) + \sum_{l=1}^L |V_l(H_{G,L}^*)| |V_l(H_{Q,L}^*)|)$.

Corollary 1. Given two graphs G and Q , the time complexity of the cross-graph learning on the CGs of G and Q (Definition 3) is no larger than that on G and Q (Definition 1).

C. Construction of CG

In this subsection, we propose an algorithm with which to construct the optimum CG of G . Since there is equivalence between the embedding of a node u at the l -th layer of GIN learning and the WL label of u at the l -th iteration of WL labeling (Sec. III-C), our idea is to use the WL label to group the nodes of G . The CG construction algorithm is shown in Algorithm 5. The time complexity of Algorithm 5 is $O(L(|V(G)| + |E(G)|))$.

Theorem 4. The CG constructed by Algorithm 5 is the optimum to minimize the time complexity in Theorem 3.

For the data graph G in \mathcal{D} , the CG of G can be precomputed. Although the CG of Q is computed on-the-fly, it is a one-off cost, as we need to perform cross-graph learning between Q and many data graphs in \mathcal{D} .

VII. EXPERIMENTAL EVALUATION

In this section, we evaluate the performance of our proposed techniques.⁴

Datasets and query workload. We conduct the experiments on three real-world datasets that are widely used for graph similarity search [3], [9], [11], [15], [55]–[57]. AIDS is an antivirus screen compound graph dataset [3], [9], [15], [55]. LINUX is a set of control-flow graphs [9], [11], [56]. PUBCHEM is a set of chemical molecule graphs [15], [55], [57].

⁴The source code is available at the project website <https://github.com/csympeng/LAN>

Following [3], [15], [55], [56], we also use a synthetic data SYN for a scalability test. SYN is generated by the generator.⁵ 20% of SYN is used by default, and 20%, 40%, 60%, 80%, and 100% of SYN are used in the scalability test. Table I shows the statistics of the datasets. Following [9], for each dataset, we sample 4,000 graphs as the query workload, which is split into training, validation and test data by 6:2:2.

For the ground-truth GED, we first compute the exact GED using the latest method [55]. If it does not finish in 10 seconds, following the approach of [9], [21], we use the best result of three approximate GED algorithms VJ [58], Hung [59], and Beam [60] as the ground-truth GED.

Metrics. The performance metrics follow the previous works [18], [19], [50]. Specifically, we use the recall at k ($recall@k$) to measure the search accuracy. $recall@k = |\mathcal{R} \cap \mathcal{R}'|/k$, where \mathcal{R} is the result set of the k -ANN search algorithm and \mathcal{R}' is the true result set. Following [18], [19], [27], [34], [50], we use *queries per second* (QPS) to measure search efficiency. QPS is the number of queries finished in a second. We focus on search performance in high recall region.

Baseline methods. We compare our method LAN with HNSW [17], which is the latest method supporting k -ANN search in a metric space and L2route [28], which is the latest learning-based routing method on PG. HNSW is used directly on graph databases. HNSW builds a hierarchy of PGs, the bottom level of which is used as the PG of LAN. Since L2route is designed for high-dimensional data, we first convert graphs into embedding vectors using Node2vec [54] and then use L2route on the embedding vectors for the k -ANN search. We use LAN_IS and HNSW_IS to denote the initial node selection methods of LAN and HNSW, respectively. Rand_IS denotes the method of randomly selecting a node as the initial node. We use LAN_Route and HNSW_Route to denote the routing methods of LAN and HNSW, respectively. We compare the cross-graph learning acceleration with HAG [46], which is the latest method that can accelerate GNN learning without reducing accuracy.

Experimental settings. The experiments are conducted using PyTorch on a server with a Quad-Core AMD Opteron CPU, 800G RAM, and a GPU card NVIDIA Tesla V100S with 32G of memory. The embedding dimensions are 128. We use the Adam optimizer. We set the initial learning rate to 0.005 and reduce it by 0.96 for every 5 epochs. We set the number of epochs to 1,000, and select the best model on validation data for testing. The neighborhood size parameter γ^* is set such that for 90% training queries, \mathcal{N}_Q can contain the 200-NNs of Q and the batch size parameter y is 20. We focus on the performance for $k = 50$. The beam size b is no smaller than k . The larger the b , the higher the recall, but the higher the query latency. Following recent works [16], [20], [50], we increase b until the target recall is achieved.

A. Comparison with existing methods

In this experiment, we compare the performance of LAN, HNSW, and L2route. LAN comprises LAN_Route, LAN_IS,

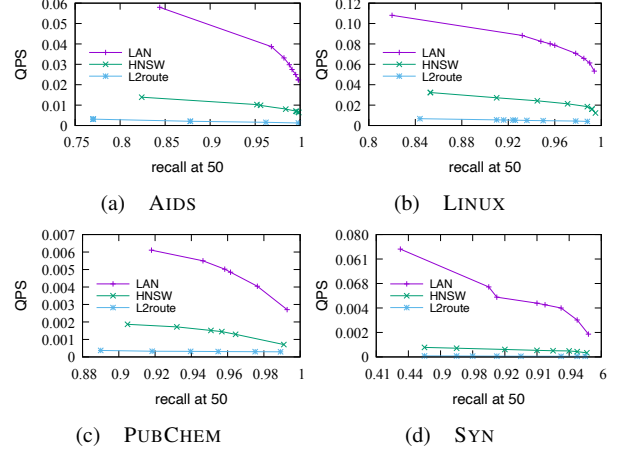


Fig. 5: Comparison with existing k -ANN search methods

and GNN acceleration. HNSW comprises HNSW_Route and HNSW_IS. Fig. 5 shows the results on AIDS, LINUX, PUBCHEM, and SYN.

In Fig. 5, we can observe that LAN significantly outperforms HNSW and L2route. In particular, at $recall@50 = 0.95$, the QPSs of LAN are $\sim 4.2x$, $\sim 3.9x$, $\sim 3.6x$, and $\sim 9x$ higher than those of HNSW on AIDS, LINUX, PUBCHEM, and SYN, respectively. At $recall@50 = 0.95$, the QPSs of LAN are $\sim 16.1x$, $\sim 18.6x$, $\sim 17x$, and $\sim 73.2x$ higher than those of L2route on AIDS, LINUX, PUBCHEM, and SYN, respectively.

In Fig. 5, we can also observe the margin of LAN decreases with the growth of recall. On the one hand, this result is reasonable, as the higher recall required, the fewer neighbors can be pruned. On the other hand, when $recall@50$ is as high as 0.98, the QPSs of LAN are still more than 3x higher than those of HNSW on AIDS, LINUX, and PUBCHEM, respectively, and more than 6x higher than that of HNSW on SYN.

We also examine GHash [9], a learning-based approximate range query method in graph databases. We run the source code of GHash and observe that on AIDS when the GED threshold is 10, the number of GED computations is 42,585, which is close to a scan of the entire dataset.

B. Routing performance with neighbor pruning

In this experiment, we compare the performances of LAN_Route and HNSW_Route. LAN comprises LAN_Route, HNSW_IS, and GNN acceleration. HNSW comprises HNSW_Route and HNSW_IS. The results are shown in Fig. 6.

In Fig. 6, we can observe that LAN_Route significantly outperforms HNSW_Route. In particular, at $recall@50 = 0.95$, the QPSs of LAN_Route are $\sim 2.9x$, $\sim 2.6x$, $\sim 2.5x$, and $\sim 5.5x$ higher than those of HNSW_Route on AIDS, LINUX, PUBCHEM, and SYN, respectively. The margin of LAN_Route with respect to HNSW_Route reduces with the growth of recall. However, when $recall@50$ is as high as 0.98, the QPSs of LAN_Route are still more than 2.5x higher than those of

⁵<https://www.cse.cuhk.edu.hk/~jcheng/graphgen1.0.zip>

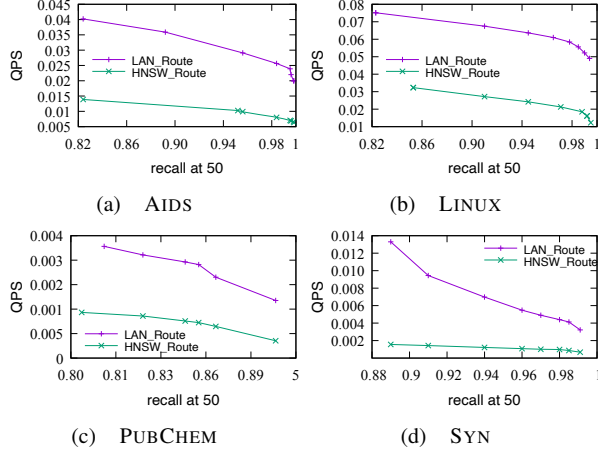


Fig. 6: Performance of routing with neighbor pruning (HNSW_IS is used for the initial node selection)

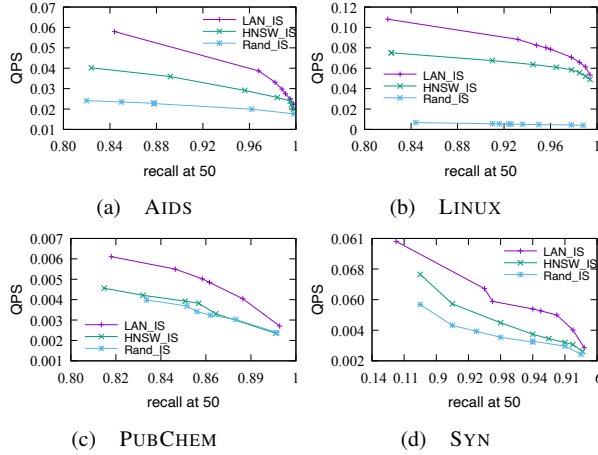


Fig. 7: Performance of initial node selection (LAN_Route is used for routing)

HNSW_Route on AIDS, LINUX, and PUBCHEM, respectively, and more than 5x higher than that of HNSW_Route on SYN.

We also examine the prediction accuracy of the neighbor ranking models during routing. The prediction accuracies are 0.73, 0.72, 0.78, and 0.7 on AIDS, LINUX, PUBCHEM, and SYN, respectively.

C. Performance of initial node selection

In this experiment, we compare the performances of LAN_IS, HNSW_IS, and Rand_IS. LAN comprises LAN_Route, LAN_IS, and GNN acceleration. HNSW comprises LAN_Route and HNSW_IS. The results are presented in Fig. 7. LAN_IS uses the optimized design of M_{nh} (Sec. V-B2) as it is faster than the basic design (Sec. V-B1).

Fig. 7 shows that LAN_IS outperforms HNSW_IS and Rand_IS. In particular, at $recall@50 = 0.95$, the QPSs of

LAN_IS are 1.4x, 1.3x, 1.3x, and 1.7x higher than those of HNSW_IS on AIDS, LINUX, PUBCHEM, and SYN, respectively. At $recall@50 = 0.95$, the QPSs of LAN_IS are $\sim 2x$, $\sim 17.2x$, $\sim 1.5x$, and $\sim 1.9x$ higher than those of Rand_IS on AIDS, LINUX, PUBCHEM, and SYN, respectively.

From Fig. 7, we can also observe that the margin of LAN_IS reduces with the growth of recall. However, at $recall@50$ is as high as 0.98, the QPSs of LAN_IS are more than 1x higher than those of HNSW_IS on AIDS, LINUX, and PUBCHEM, respectively, and more than 1.5x higher than that of HNSW_IS on SYN. At $recall@50 = 0.98$, the QPSs of LAN_IS are 1.5x, 15.1x, 1.3x, and 1.7x higher than those of Rand_IS on AIDS, LINUX, PUBCHEM, and SYN, respectively.

The precision of our initial node prediction model is shown in Fig. 8, where we can observe that the precisions exceed 0.7 on AIDS, LINUX, PUBCHEM, and SYN, respectively. Since $1 - (1 - 0.7)^4 > 0.99$, we only need to randomly sample 4 nodes from \hat{N}_Q , such that the probability that at least one sample is in N_Q is larger than 0.99.

We also compare the basic and optimized designs of the neighborhood prediction model. The time costs of prediction using the optimized design are only 3.3%, 4.1%, 3.6%, and 2.2% of the basic design on AIDS, LINUX, PUBCHEM, and SYN, respectively.




Fig. 8: Accuracy of initial node prediction

D. Scalability evaluation

In this experiment, we evaluate the scalability of LAN. We focus on the SYN dataset. Fig. 9 presents the results. In Fig. 9, the x-axis is the scale of the dataset and the y-axis is the average running time of a query.

From Fig. 9, we can observe that LAN scales linearly with the size of the dataset. The reason is that following the approach of [18], [19], [50] to support a large dataset, we randomly split the dataset into equal-size sub-datasets and sequentially perform k -ANN search on each sub-dataset. The top- k in the search results on all sub-datasets are returned as the final result of the whole dataset. From Fig. 9, we can also observe that the gap between $recall@50 = 0.9$ and $recall@50 = 0.95$ is much smaller than the gap between $recall@50 = 0.95$ and $recall@50 = 0.98$. This finding is consistent with the observations in previous experiments that the increase in running time is more than a linear function of the increase in recall.

E. Performance of cross-graph learning acceleration

Fig. 10 presents the effects of cross-graph learning acceleration on the efficiency of k -ANN search. From Fig. 10, we can

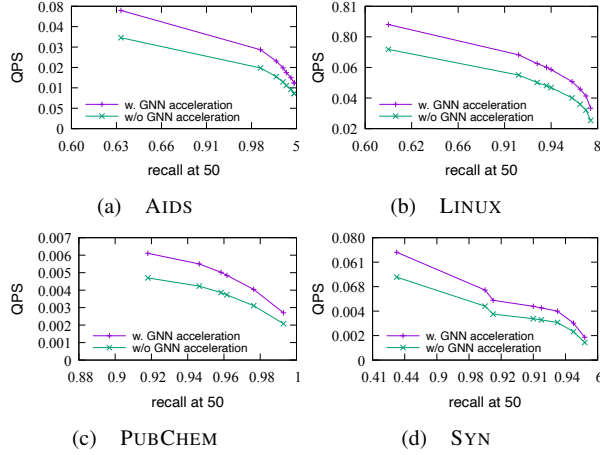


Fig. 10: Performance of cross-graph learning acceleration

observe that the QPS is increased after using our cross-graph learning acceleration technique. In particular, at $\text{recall}@50 = 0.95$, the QPSs are increased by $\sim 17\%$, $\sim 18\%$, $\sim 15\%$, and $\sim 17\%$ on AIDS, LINUX, PUBCHEM, and SYN, respectively.

Fig. 11 shows the breakdown of the query time before using our cross-graph learning acceleration technique. It is evident that cross-graph learning accounts for $\sim 24\%$, $\sim 25\%$, $\sim 20\%$, and $\sim 29\%$ of the query time on AIDS, LINUX, PUBCHEM, and SYN, respectively.

Fig. 12 shows the extent to which cross-graph learning is speeded up using our acceleration technique. The speedup is based on the original cross-graph learning technique. We can see that cross-graph learning is speeded up by $\sim 4x$, $\sim 4.2x$, $\sim 5.3x$, and $\sim 3.1x$ on AIDS, LINUX, PUBCHEM, and SYN, respectively. In contrast, HAG cannot speedup the cross-graph learning.

We remark that the GNN is only used to predict which graphs should be pruned from GED computation; we need to compute GEDs for the graphs that are not pruned. GED computation dominates the overall query time, which can be observed in the performance breakdown (Fig. 11). Therefore, although the GNN on PUBCHEM has the largest speedup, the improvement of QPS on PUBCHEM is not the largest.




Fig. 11: Breakdown of k -ANN search time

F. Index construction time

The PG of LAN is the bottom level PG in the PG hierarchy of HNSW. Therefore, we use the time of HNSW for constructing

the bottom level PG as the index construction time of LAN. We use 5 threads to process multiple GED computation tasks concurrently. The index construction times of LAN on AIDS, LINUX, PUBCHEM, and SYN are 6.8, 77.5, 120.2, and 7.8 hours, respectively.

VIII. CONCLUSION

In this paper, we propose a learning-based k -ANN search method in graph databases. The core of our method uses graph learning models to avoid unnecessary distance computations during the routing on PG. First, we propose a routing algorithm with neighbor pruning. At each routing step, a graph learning model is used to rank the neighbors and only the top neighbors are explored. Second, we propose a graph learning model for initial node selection, such that the initial node selected has a high probability of being in the neighborhood of Q . Third, we propose a compressed GNN-graph to accelerate the cross-graph learning used in the neighbor ranking and initial node selection. The efficiency of learning is improved without degrading its accuracy. Our experiments show that our method is effective and significantly outperforms the state-of-the-art k -ANN search methods on real-world benchmark graph datasets.

In the future, we plan to study distributed k -ANN search methods to support larger graph databases.

APPENDIX

This appendix contains the proofs of the theorems and lemmas. Some detailed derivations are available on the project website.

A. Proof of Lemma 1

To prove Lemma 1, we first prove the following two Lemmas.

Lemma 3. For any routing step of `np_route`, let S denote the set of explored nodes so far and $N_G(S)$ denote the neighbors of the nodes in S in \mathcal{G} . The node G explored at this step must be the node that is the closest to Q among all unexplored nodes in $N_G(S)$.

Proof. It is trivial for the first stage of routing (*i.e.*, the routing steps from the initial node to the first local optimal; Lines 1-12 of Algorithm 4). The reason is as follows. At each routing step, suppose G is the current node. All neighbors of G with distances to Q no more than $d(Q, G)$ are added to W , and the node in W that is the closest to Q is explored at the next routing step.

For the second stage (*i.e.*, the routing steps from the first local optimal to the end of the routing), the while-loop (Lines 22-27) performs the routing steps. Suppose the distance threshold for the current routing step is γ (Line 22). Let $N_G^{ue}(S, \gamma)$ denote the unexplored nodes in $N_G(S)$ with distances to Q that are no more than γ . Lines 15-17 ensure that all nodes in $N_G^{ue}(S, \gamma)$ are added to W . Line 18 may squeeze some nodes in $N_G^{ue}(S, \gamma)$ out of W . If the node G in $N_G^{ue}(S, \gamma)$ that is the closest to Q is squeezed out, the routing stops. Otherwise, G

must be the first to be explored in the while-loop (Lines 22-27). \square

Lemma 4. For any routing step of `baseline`, let S denote the set of explored nodes so far and $N_G(S)$ denote the neighbors of the nodes in S in G . The node G explored at this step must be the node that is the closest to Q among all unexplored nodes in $N_G(S)$.

Proof. It is trivial as, at each routing step, let G be the current node, all neighbors of G are added to W , and the node in W that is the closest to Q is explored in the next routing step. \square

The proof sketch of Lemma 1 is as follows.

Proof. Let seq_{bs}^i and seq_{np}^i denote the sequences of explored nodes of `baseline` and `np_route` from the 0-th step to the i -th step, respectively. We prove this lemma by induction.

For the 0-th step, it is trivial as $seq_{bs}^0 = seq_{np}^0 = [G_{init}]$.

Suppose $seq_{bs}^{i-1} = seq_{np}^{i-1}$. Due to Lemma 3 and Lemma 4, the nodes explored by `baseline` and `np_route` at the i -th step must be the same. Therefore, $seq_{bs}^i = seq_{np}^i$. \square

B. Proof of Theorem 1

Proof. Since the sequences of nodes explored by `np_route` and `baseline` are the same, the search results \mathcal{R}_{np} returned by `np_route` must be the same as the search result \mathcal{R}_{bs} of `baseline`.

Since `np_route` prunes some neighbors, the number of distance computations must be no more than that of `baseline`.

The space cost of `baseline` is the linear to $|S| + |N_G(S)|$, where S denotes the explored nodes and $N_G(S)$ denotes the neighbors of the nodes S in G . Since `np_route` needs to record the batches, the space cost of `np_route` is higher than that of `baseline`. Extremely, the size of each batch is 1, where space cost of `np_route` is $|S| + 2 \times |N_G(S)|$, as the explored nodes of `np_route` are the same as that of `baseline` (see Lemma 1). Therefore, the space cost of `np_route` is at most two times of that of `baseline`. \square

C. Proof of Lemma 2

Proof. If we randomly sample a graph G in \mathcal{N}_Q , G is in \mathcal{N}_Q with a probability of p . Since the s samples are sampled independently, the probability that all the s samples are not in \mathcal{N}_Q is $(1-p)^s$. Therefore, the probability that at least one sample is in \mathcal{N}_Q is $1 - (1-p)^s$. \square

D. Proof of Theorem 2

Proof. To prove $\mathbf{h}_{H_{G,L}^*} || \mathbf{h}_{H_{Q,L}^*} = \mathbf{h}_G || \mathbf{h}_Q$, we only need to prove $\mathbf{h}_{H_{G,L}^*} = \mathbf{h}_G$ and $\mathbf{h}_{H_{Q,L}^*} = \mathbf{h}_Q$. In the following, we only prove that $\mathbf{h}_{H_{G,L}^*} = \mathbf{h}_G$ as $\mathbf{h}_{H_{Q,L}^*} = \mathbf{h}_Q$ can be proved in the same way.

In proving $\mathbf{h}_{H_{G,L}^*} = \mathbf{h}_G$, the key is to prove that in each layer of cross-graph learning, if a node is in a group, the embedding of the node computed by the cross-graph learning on G and Q equals to the embedding of the group computed by the cross-graph learning on the CG of G and the CG of Q .

In the l -th layer of the cross-graph learning on G and Q , for a node u of G , u needs to i) compute \mathbf{t}_u^l (i.e., the aggregate embeddings from its neighbors in G), ii) compute μ_u^l (i.e., the aggregate weighted embeddings from all nodes in Q), and iii) multiply \mathbf{W}^l .

In the l -th layer of the cross-graph learning on the CG of G and the CG of Q , a node g at the l -th level of the CG of G needs to i) compute \mathbf{t}_g^l (i.e., the aggregate weighted embeddings from its incoming neighbors in the $l-1$ -th level of the CG of G), ii) compute μ_g^l (i.e., the aggregate weighted embeddings from all nodes at the l -th level of the CG of Q), and iii) multiply \mathbf{W}^l .

By Definition 3, if the node u is in the group g , $\mathbf{t}_u^l = \mathbf{t}_g^l$ and $\mu_u^l = \mu_g^l$. Therefore, $\mathbf{h}_u^l = \mathbf{h}_g$.

In the L -th layer, since $\mathbf{h}_G = \text{mean}_{u \in G} \mathbf{h}_u^L$, $\mathbf{h}_{H_{G,L}^*}$ is the weighted average of the embeddings of the nodes in the L -th level of $H_{G,L}^*$, where the weight is group size, and each node u in G must belong to a node in the L -th level of $H_{G,L}^*$, $\mathbf{h}_G = \mathbf{h}_{H_{G,L}^*}$ must hold. \square

E. Proof of Theorem 3

Proof. The cross-graph learning at the l -th level of $H_{G,L}^*$, $l = 1, \dots, L$, has three steps: i) collect information from the incoming neighbors in the $l-1$ -th level of $H_{G,L}^*$, ii) compute attention for the nodes in the l -th level of $H_{Q,L}^*$, and iii) multiply \mathbf{W} . It is the same for the l -th level of $H_{Q,L}^*$.

For i), the time cost is linear to the number of incoming edges of the l -th level. The total time for L levels is $O(|E(H_{G,L}^*)| + |E(H_{Q,L}^*)|)$.

For ii), the time cost is linear to the number of pairs of the nodes in the l -th level of $H_{G,L}^*$ and the nodes in the l -th level of $H_{Q,L}^*$. The total time for L levels is $O(\sum_{l=1}^L |V_l(H_{G,L}^*)| |V_l(H_{Q,L}^*)|)$. Note that we do not need to compute the attention for the 0-th level.

For iii), the time cost is linear to the number of nodes. Hence, the total time is $O(|V(H_{Q,L}^*)| + |V(H_{G,L}^*)|)$. \square

F. Proof of Corollary 1

Proof. It is trivial as, for a graph G , the number of edges between two levels of the CG of G is no larger than $E(G)$ and the number of nodes in each level of the CG of G is no larger than $V(G)$. \square

G. Proof of Theorem 4

Proof. The key argument here is that the WL label of a node is equivalent to the embedding of the node. If the nodes have been grouped by their WL labels, it is impossible to further reduce the number of groups while making sure that all nodes in a group always have the same embedding. \square

ACKNOWLEDGMENT

Thanks to the support of HKRGC GRF 12201119 and 12201518, CRF C6030-18GF, IRCMS/19-20/H01, NSF of China for Joint Fund Project U1936218, Guangdong Basic and Applied Basic Research Foundation 2019B1515130001, and NSF of Shandong Province ZR2019LZH006.

REFERENCES

- [1] A. A. Collado, R. Carrasco-Velaz, N. García-Pedrajas, and G. Cerruela-García, "Maximum common property: A new approach for molecular similarity," *Journal of Cheminformatics*, vol. 12, no. 61, 2020.
- [2] H. Du, "Chemical molecules search based on graph similarity measure," in *Proc. of the International Conference on Biomedical Engineering and Computer Science*, 2010, pp. 1–4.
- [3] X. Zhao, C. Xiao, X. Lin, W. Zhang, and Y. Wang, "Efficient structure similarity searches: a partition-based approach," *The VLDB Journal*, vol. 27, no. 1, pp. 53–78, 2018.
- [4] C. García-Hernández, A. Fernández, and F. Serratos, "Ligand-based virtual screening using graph edit distance as molecular similarity measure," *Journal of Chemical Information and Modeling*, vol. 59, no. 4, pp. 1410–1421, 2019.
- [5] M. Naderi, C. Alvin, Y. Ding, S. Mukhopadhyay, and M. Brylinski, "A graph-based approach to construct target-focused libraries for virtual screening," *Journal of cheminformatics*, vol. 8, no. 1, pp. 1–16, 2016.
- [6] Y. Zhu, L. Qin, J. X. Yu, and H. Cheng, "Answering top-*k* graph similarity queries in graph databases," *IEEE Transactions on Knowledge and Data Engineering*, vol. 32, no. 8, pp. 1459–1474, 2019.
- [7] R. Wang, T. Zhang, T. Yu, J. Yan, and X. Yang, "Combinatorial learning of graph edit distance via dynamic embedding," in *Proc. of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2021, pp. 5241–5250.
- [8] Z. Abu-Aisheh, R. Raveaux, J.-Y. Ramel, and P. Martineau, "An exact graph edit distance algorithm for solving pattern recognition problems," in *Proc. of the International Conference on Pattern Recognition Applications and Methods*, 2015, pp. 271–278.
- [9] Z. Qin, Y. Bai, and Y. Sun, "GHashing: Semantic graph hashing for approximate similarity search in graph databases," in *Proc. of the ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2020, pp. 2062–2072.
- [10] S. Cesare and Y. Xiang, "Malware variant detection using similarity search over sets of control flow graphs," in *Proc. of the International Conference on Trust, Security and Privacy in Computing and Communications*, 2011, pp. 181–189.
- [11] X. Wang, X. Ding, A. K. Tung, S. Ying, and H. Jin, "An efficient graph indexing method," in *Proc. of the IEEE International Conference on Data Engineering*, 2012, pp. 210–221.
- [12] Y. Zhu, L. Qin, J. X. Yu, and H. Cheng, "Finding top-*k* similar graphs in graph databases," in *Proc. of the International Conference on Extending Database Technology*, 2012, pp. 456–467.
- [13] Z. Abu-Aisheh, R. Raveaux, and J.-Y. Ramel, "Efficient *k*-nearest neighbors search in graph space," *Pattern Recognition Letters*, vol. 134, pp. 77–86, 2020.
- [14] Z. Zeng, A. K. H. Tung, J. Wang, J. Feng, and L. Zhou, "Comparing stars: On approximating graph edit distance," *Proc. VLDB Endow.*, vol. 2, no. 1, pp. 25–36, 2009.
- [15] J. Kim, "Boosting graph similarity search through pre-computation," in *Proc. of the ACM SIGMOD International Conference on Management of Data*, 2021, pp. 951–963.
- [16] C. Li, M. Zhang, D. G. Andersen, and Y. He, "Improving approximate nearest neighbor search through learned adaptive early termination," in *Proc. of the ACM SIGMOD International Conference on Management of Data*, 2020, pp. 2539–2554.
- [17] Y. A. Malkov and D. A. Yashunin, "Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 42, no. 4, pp. 824–836, 2020.
- [18] C. Fu, C. Xiang, C. Wang, and D. Cai, "Fast approximate nearest neighbor search with the navigating spreading-out graph," vol. 12, no. 5, pp. 461–474, 2019.
- [19] C. Fu, C. Wang, and D. Cai, "High dimensional similarity search with satellite system graph: Efficiency, scalability, and unindexed query compatibility," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, pp. 1–1, 2021.
- [20] W. Li, Y. Zhang, Y. Sun, W. Wang, M. Li, W. Zhang, and X. Lin, "Approximate nearest neighbor search on high dimensional data — experiments, analyses, and improvement," *IEEE Transactions on Knowledge and Data Engineering*, vol. 32, no. 8, pp. 1475–1488, 2020.
- [21] L. Yang and L. Zou, "Noah: Neural-optimized A* search algorithm for graph edit distance computation," in *Proc. of IEEE International Conference on Data Engineering (ICDE)*, 2021.
- [22] Y. Li, C. Gu, T. Dullien, O. Vinyals, and P. Kohli, "Graph matching networks for learning the similarity of graph structured objects," in *Proc. of International Conference on Machine Learning (ICML)*, 2019, pp. 3835–3845.
- [23] B. Harwood and T. Drummond, "Fannng: Fast approximate nearest neighbour graphs," in *Proc. of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 5713–5722.
- [24] L. C. Shimomura, R. S. Oyamada, M. R. Vieira, and D. S. Kaster, "A survey on graph-based methods for similarity searches in metric spaces," *Information Systems*, vol. 95, p. 101507, 2021.
- [25] J. Qin, W. Wang, C. Xiao, and Y. Zhang, "Similarity query processing for high-dimensional data," *Proc. VLDB Endow.*, vol. 13, no. 12, pp. 3437–3440, 2020.
- [26] L. Prokhorenkova and A. Shekhovtsov, "Graph-based nearest neighbor search: From practice to theory," in *Proc. of the International Conference on Machine Learning*, vol. 119, 2020, pp. 7803–7813.
- [27] W. Zhao, S. Tan, and P. Li, "Song: Approximate nearest neighbor search on GPU," in *Proc. of the IEEE International Conference on Data Engineering (ICDE)*, 2020, pp. 1033–1044.
- [28] D. Baranchuk, D. Persiyanov, A. Sinitin, and A. Babenko, "Learning to route in similarity graphs," in *Proc. of the International Conference on Machine Learning*, 2019, pp. 475–484.
- [29] P. Lin and W.-L. Zhao, "A comparative study on hierarchical navigable small world graphs," *ArXiv*, vol. abs/1904.02077, 2019.
- [30] J. Vargas Muñoz, M. A. Gonçalves, Z. Dias, and R. da S. Torres, "Hierarchical clustering-based graphs for large scale approximate nearest neighbor search," *Pattern Recognition (PR)*, vol. 96, p. 106970, 2019.
- [31] Y. Malkov, A. Ponomarenko, A. Logvinov, and V. Krylov, "Approximate nearest neighbor algorithm based on navigable small world graphs," *Information Systems*, vol. 45, pp. 61–68, 2014.
- [32] K. Aoyama, A. Ogawa, T. Hattori, and T. Hori, "Double-layer neighborhood graph based similarity search for fast query-by-example spoken term detection," in *Proc. of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2015, pp. 5216–5220.
- [33] K. Lu, H. Wang, W. Wang, and M. Kudo, "VHP: Approximate nearest neighbor search via virtual hypersphere partitioning," *Proc. VLDB Endow.*, vol. 13, no. 9, pp. 1443–1455, 2020.
- [34] L. Prokhorenkova and A. Shekhovtsov, "Graph-based nearest neighbor search: From practice to theory," in *Proc. of the International Conference on Machine Learning*, 2020, pp. 7803–7813.
- [35] H. Wang, Z. Wang, W. Wang, Y. Xiao, Z. Zhao, and K. Yang, "A note on graph-based nearest neighbor search," *CoRR*, vol. abs/2012.11083, 2020.
- [36] Y. Bai, H. Ding, K. Gu, Y. Sun, and W. Wang, "Learning-based efficient graph similarity computation via multi-scale convolutional set matching," in *Proc. of AAAI Conference on Artificial Intelligence (AAAI)*, 2020, pp. 3219–3226.
- [37] Y. Peng, B. Choi, and J. Xu, "Graph edit distance learning via modeling optimum matchings with constraints," in *Proc. of the International Joint Conference on Artificial Intelligence (IJCAI)*, 2021, pp. 1534–1540.
- [38] —, "Graph learning for combinatorial optimization: A survey of state-of-the-art," *Data Sci. Eng.*, vol. 6, no. 2, pp. 119–141, 2021.
- [39] W. L. Hamilton, R. Ying, and J. Leskovec, "Inductive representation learning on large graphs," in *Proc. of the International Conference on Neural Information Processing Systems*, 2017, pp. 1025–1035.
- [40] J. Chen, T. Ma, and C. Xiao, "FastGCN: Fast learning with graph convolutional networks via importance sampling," in *Proc. of the International Conference on Learning Representations*, 2018.
- [41] W. Huang, T. Zhang, Y. Rong, and J. Huang, "Adaptive sampling towards fast graph representation learning," in *Proc. of the International Conference on Neural Information Processing Systems*, 2018, pp. 4563–4572.
- [42] F. Wu, A. Souza, T. Zhang, C. Fifty, T. Yu, and K. Weinberger, "Simplifying graph convolutional networks," in *Proc. of the International Conference on Machine Learning*, vol. 97, 2019, pp. 6861–6871.
- [43] Y. Ye and S. Ji, "Sparse graph attention networks," *IEEE Transactions on Knowledge and Data Engineering*, pp. 1–1, 2021.
- [44] S. A. Tailor, J. Fernandez-Marques, and N. D. Lane, "Degree-quant: Quantization-aware training for graph neural networks," in *Proc. of the International Conference on Learning Representations*, 2021.
- [45] W. Wu, B. Li, C. Luo, and W. Nejdl, "Hashing-accelerated graph neural networks for link prediction," in *Proceedings of the Web Conference 2021*, 2021, pp. 2910–2920.

- [46] Z. Jia, S. Lin, R. Ying, J. You, J. Leskovec, and A. Aiken, "Redundancy-free computation for graph neural networks," in *Proc. of the ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2020, pp. 997–1005.
- [47] H. Liu, S. Lu, X. Chen, and B. He, "G3: When graph neural networks meet parallel graph processing systems on GPUs," *Proc. of the VLDB Endowment*, vol. 13, no. 12, pp. 2813–2816, 2020.
- [48] Z. Chen, M. Yan, M. Zhu, L. Deng, G. Li, S. Li, and Y. Xie, "fuseGNN: Accelerating graph convolutional neural network training on GPGPU," in *Proc. of the IEEE/ACM International Conference On Computer Aided Design (ICCAD)*, 2020, pp. 1–9.
- [49] H. Bunke, "On a relation between graph edit distance and maximum common subgraph," *Pattern Recognition Letters*, vol. 18, no. 8, pp. 689–694, 1997.
- [50] M. Wang, X. Xu, Q. Yue, and Y. Wang, "A comprehensive survey and experimental comparison of graph-based approximate nearest neighbor search," *Proc. VLDB Endow.*, vol. 14, no. 11, pp. 1964–1978, 2021.
- [51] K. Xu, W. Hu, J. Leskovec, and S. Jegelka, "How powerful are graph neural networks?" in *Proc. of the International Conference on Learning Representations, ICLR*, 2019.
- [52] X. He, J. Pan, O. Jin, T. Xu, B. Liu, T. Xu, Y. Shi, A. Atallah, R. Herbrich, S. Bowers, and J. Q. n. Candela, "Practical lessons from predicting clicks on ads at Facebook," in *Proc. of the International Workshop on Data Mining for Online Advertising*, 2014, pp. 1–9.
- [53] C.-Y. Chiu, A. Prayoonwong, and Y.-C. Liao, "Learning to index for nearest neighbor search," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 42, no. 8, pp. 1942–1956, 2020.
- [54] A. Grover and J. Leskovec, "Node2vec: Scalable feature learning for networks," in *Proc. of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 855–864.
- [55] L. Chang, X. Feng, X. Lin, L. Qin, W. Zhang, and D. Ouyang, "Speeding up GED verification for graph similarity search," in *Proc. of IEEE International Conference on Data Engineering (ICDE)*, 2020, pp. 793–804.
- [56] K. Gouda and M. Hassaan, "CSI_GED: An efficient approach for graph edit similarity computation," in *Proc. of the IEEE International Conference on Data Engineering (ICDE)*, 2016, pp. 265–276.
- [57] Y. Peng, Z. Fan, B. Choi, J. Xu, and S. S. Bhowmick, "Authenticated subgraph similarity search in outsourced graph databases," *IEEE Transactions on Knowledge and Data Engineering*, vol. 27, no. 7, pp. 1838–1860, 2015.
- [58] S. Fankhauser, K. Riesen, and H. Bunke, "Speeding up graph edit distance computation through fast bipartite matching," in *Proc. of International Workshop on Graph-Based Representations in Pattern Recognition*, 2011, pp. 102–111.
- [59] K. Riesen and H. Bunke, "Approximate graph edit distance computation by means of bipartite graph matching," *Image and Vision computing*, vol. 27, no. 7, pp. 950–959, 2009.
- [60] M. Neuhaus, K. Riesen, and H. Bunke, "Fast suboptimal algorithms for the computation of graph edit distance," in *Structural, Syntactic, and Statistical Pattern Recognition*, 2006, pp. 163–172.